

IE 3301 - 004
ENGINEERING PROBABILITY

REAL-WORLD DATA: ANALYSIS

Part I

November 20, 2017

Joe Cloud
Student ID: 1000921236
University of Texas at Arlington

*I, Joe Cloud, did not give or receive any assistance on
this project, and the report submitted is wholly my own.*

Contents

| | |
|--|----|
| Introduction | 2 |
| Data | 2 |
| Set One | 2 |
| Set Two | 2 |
| Descriptive Analysis | 3 |
| Set One | 3 |
| Set Two | 4 |
| Appendix | 5 |
| I - Dataset: One | 5 |
| II - Dataset: Two | 5 |
| III - Tables | 5 |
| A - Set One: Frequency table | 5 |
| B - Set Two: Frequency table | 6 |
| IV - Code | 6 |
| A - Descriptive Analysis | 6 |
| B - Data Processing: Set One | 8 |
| C - Data Processing: Set Two | 8 |
| V - Equipment Used | 10 |
| References | 11 |

INTRODUCTION

The aim of this project, as stated in the project brief is to analyze real-world data using techniques covered in the introductory Engineering Probability course. For this portion of the project, data collection was performed on two independent data sets: one is a sample of a continuous random variable that is suspected to be normally distributed, while the second is a measure of the time interval between log events; then each dataset is statistically summarized. For my project, I chose to collect a set of resistor values from a batch of $100\ 1\text{k}\Omega \pm 5\%$ resistors ($k\Omega$ is the measurement of resistance, in units of Kilo Ohms), for the first set. For the second set, I compiled login logs for users from a compute cluster and ran preprocessing scripts to extrace the login intervals. The result was a dataset of 120 events containing the login interval (in seconds) between user logins.

DATA

Set One

Data collection began with the process of measuring the resistance of each resistor and recording its value. The measurements were conducted with a calibrated 4.5 digit digital multimeter (see Appendix V for equipment information). To perform this measurement, a digital multimeter was placed in 'resistance' measuring mode, and with probes attached to the correct ports, each resistor was attached to the end of the probe on either side of the filament. Once attached the resistive value would be flashed on the display of the digital multimeter. Each measurement was entered into a spreadsheet and later outputted to a CSV file for further processing by data analysis scripts.

Before data analysis was performed, 0.33Ω was subtracted from each value, this is done to account for the resistance added by the multimeter probes. Although this minor offset did not impact the distribution of measurements for the purposes of the project, it did provide us with results that more accurately represent the true values. The values are offset by approximately 0.03%.

In an ideal world, the resistors would measure identically to the $1\text{k}\Omega$ label. In reality, the cost to refine manufacturing to attain such level of accuracy is expensive. Tolerances guarantee a range of possible [random] values which are acceptable for most electronics. This dataset was created to explore the distribution of these values.

Set Two

The data collection for Set Two revolved around login access to a compute cluster and will be difficult to replicate without access to a system with comprehensive logging and login activity. Though, it is possible for the reader to perform similar analysis with data from a personal machine's logged data. On a Unix-like operating system this can be performed with the use of the `last` command. One of the benefits of computer-triggered collection is that it is much less susceptible to timing errors as opposed to that of a 'human polling' based approach.

I included the source code necessary to collect the data in Appendix IV.C. The process is largely computational. After the raw data was collected with the `last` command, the data was parsed through a preprocessing script written to convert the timestamps in HH:MM:SS format to that of total seconds. Then, each event was subtracted from the next, yielding interlogin times. The negative sign is discarded as only the magnitude is significant. As a side note, all source code written for this project are included in the appendix.

DESCRIPTIVE ANALYSIS

This section details the results of analysis performed on both datasets One and Two. To analyze the data, python scripts were written to perform the requested calculations and output the desired plots for each dataset. The scripts are included as part of the Code section of Appendix IV.A.

Set One

The first step in analysis for Set One was calculating the mean \bar{X}_1 , which is $0.98349 \text{ k}\Omega$, between a minimum value of $0.97987 \text{ k}\Omega$ and a maximum value of $0.98827 \text{ k}\Omega$. The next step was calculating the standard deviation, σ_1 . This was done using `numpy` as opposed to doing it by hand with $\sigma_1 = \sqrt{\frac{1}{N-1} \sum_{i=1}^N (X_i - \bar{X}_1)^2}$. The result for σ_1 was 0.001739 , which would indicate that the spread from the mean is fairly low, considering the significance of the value. The last calculation necessary for box-and-whisker plot was the quartiles, in which $Q1 = 0.98257 \text{ k}\Omega$, $Q2 = 0.98659 \text{ k}\Omega$, $Q3 = 0.98467 \text{ k}\Omega$.

The box-and-whisker plot is a helpful method of visualizing spread within a dataset. The vertical line segments end on either side of the box represent the minimum and maximum values within the sample space, while the dot to the right represents outliers within the data. The box itself is useful for showing the quartiles (first being the left edge of the box, with third being the right). In this set, the data appears to be dispersed somewhat evenly. The whiskers are similar in length, with only a slight bias towards the left. This would indicate that the resistant of the $1 \text{ k}\Omega$ hold an average lower than the median resistance. (x-axis is resistance in $\text{k}\Omega$)

The frequency table (Appendix III.A) was created first by determining a bin size. The range of the data was divided into 10 equal parts. Each bin would contain the tally of the values within a particular slice. For the resistor values, which is a continuous random variable that we suspect to be normally distributed, it is promising to see that the tallies peak where the center-most values are located. The right-most extreme values do appear to decrease quite significantly, while the lower end appears to contain more tallies than we would expect for a proper normal distribution.

The histogram represents the distribution of the resistive values across 10 bins, where the y-axis is the number of values tallied within the slice, and x-axis is the measured resistance in units of $\text{k}\Omega$. Based on the




Figure 1: Set One: Box-and-whisker plot




Figure 2: Set One: Histogram

histogram and previous data analysis, the sample of resistors do not follow a normal distribution due to a lack of symmetry.

Set Two

The first step in analysis for Set Two was calculating the mean \bar{X}_2 , which is 8445.05 seconds (s), between a minimum value of 1.0s (restricted by measurement resolution) and a maximum value of 83749s. The next step was calculating the standard deviation, σ_2 , which was 15392.12s. This indicates that spread from the mean is high. The last calculation necessary for box-and-whisker plot was the quartiles, in which $Q1 = 251.75s$, $Q2 = 1865.0s$, $Q3 = 8052.25s$.

In this set, the data appears to be heavily right skewed, with the majority of the values located on the lower end of the scale (in seconds) and a few large outliers to the right. This foreshadows an interesting frequency distribution and histogram. Since this interval data is extracted from timestamps, we suspect the data to follow an exponential distribution. It is based on a continuous random variable (through the time stamp), but the values are discretized due to the limits of logging precision.

Similar to Set One, the frequency table (Appendix III.B) was created first by determining a bin size. The range of the data was divided into 10 equal parts as well. Each bin contained the tally of the values within a particular slice. When tabled, the values for the login intervals showed that the vast majority of tallies were entered into the first slice, and a significant decrease in each successive step. The second through the last bin combined makes up less than a third of the tallies of the first bin.

The histogram to the right represents the distribution of the login interval values across 10 bins. Where the y-axis is the number of values tallied within the slice, and x-axis is time between successive logins in seconds. Based on the histogram and previous data analysis, I conclude that the data follows an exponential distribution. As interval increases there is an exponential decrease in the values recorded, with some outliers to the extreme but with more samples the overall distribution would continue to show strong resemblance to that of the standard exponential distribution.




Figure 3: Set Two: Box-and-whisker plot




Figure 4: Set Two: Histogram

APPENDIX

I - Dataset: One

Measured resistance of each resistor in the batch of 100.

| | | | | | | | | | |
|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|
| 0.98227 | 0.98467 | 0.98377 | 0.98277 | 0.97987 | 0.98467 | 0.98437 | 0.98337 | 0.98317 | 0.98507 |
| 0.98027 | 0.98077 | 0.98487 | 0.98037 | 0.98457 | 0.98237 | 0.98387 | 0.98367 | 0.98257 | 0.98487 |
| 0.98377 | 0.98217 | 0.98017 | 0.98487 | 0.98447 | 0.98737 | 0.98317 | 0.98287 | 0.98477 | 0.98417 |
| 0.98207 | 0.98387 | 0.98757 | 0.98427 | 0.98477 | 0.98257 | 0.98537 | 0.98517 | 0.98037 | 0.98007 |
| 0.98527 | 0.98617 | 0.98397 | 0.98127 | 0.98357 | 0.98367 | 0.98387 | 0.98097 | 0.98357 | 0.98077 |
| 0.98287 | 0.98087 | 0.98627 | 0.98107 | 0.98377 | 0.98327 | 0.98537 | 0.98357 | 0.98577 | 0.98547 |
| 0.98247 | 0.98507 | 0.98337 | 0.98367 | 0.98527 | 0.98347 | 0.98827 | 0.98207 | 0.98337 | 0.98297 |
| 0.98297 | 0.98097 | 0.98387 | 0.98287 | 0.98467 | 0.98327 | 0.98157 | 0.98037 | 0.98487 | 0.98117 |
| 0.98547 | 0.98397 | 0.98437 | 0.98337 | 0.98317 | 0.98287 | 0.98507 | 0.98397 | 0.98287 | 0.98327 |
| 0.98167 | 0.98197 | 0.98317 | 0.98547 | 0.98377 | 0.98637 | 0.98057 | 0.98277 | 0.98547 | 0.98447 |

II - Dataset: Two

The interlogin time for users on the compute cluster, in seconds. 120 interlogin times recorded.

| | | | | | | | | | | | |
|------|-------|-------|-------|-------|-------|-------|-------|-------|-------|------|-------|
| 1796 | 67183 | 7901 | 13437 | 6645 | 916 | 41159 | 8269 | 2684 | 38237 | 3263 | 47470 |
| 6397 | 24371 | 28258 | 12257 | 1348 | 18621 | 6675 | 2691 | 25385 | 86 | 10 | 18647 |
| 68 | 31 | 19 | 2202 | 124 | 508 | 3 | 72894 | 5630 | 15758 | 7980 | 431 |
| 3918 | 3541 | 11982 | 35755 | 6654 | 2526 | 15853 | 3533 | 507 | 367 | 1404 | 2306 |
| 2237 | 23714 | 887 | 141 | 5493 | 520 | 83749 | 3516 | 29573 | 27200 | 2 | 32295 |
| 2197 | 3579 | 4556 | 166 | 864 | 56 | 113 | 98 | 5 | 1645 | 3 | 7368 |
| 1846 | 2136 | 2214 | 213 | 535 | 632 | 1440 | 268 | 243 | 303 | 1182 | 54091 |
| 1 | 19565 | 690 | 243 | 83 | 135 | 84 | 3346 | 6641 | 1373 | 457 | 393 |
| 19 | 187 | 376 | 9292 | 1214 | 589 | 567 | 1240 | 1884 | 207 | 423 | 5 |
| 905 | 44505 | 71 | 11948 | 10404 | 2151 | 2413 | 227 | 10287 | 196 | 25 | 8480 |

III - Tables

A - Set One: Frequency table

| Value Range | Count |
|----------------------------|-------|
| $0.97987 \leq X < 0.98071$ | 8 |
| $0.98071 \leq X < 0.98155$ | 8 |
| $0.98155 \leq X < 0.98239$ | 8 |
| $0.98239 \leq X < 0.98407$ | 16 |
| $0.98407 \leq X < 0.98491$ | 25 |
| $0.98491 \leq X < 0.98575$ | 17 |
| $0.98575 \leq X < 0.98659$ | 12 |
| $0.98659 \leq X < 0.98743$ | 4 |
| $0.98743 \leq X < 0.98827$ | 1 |
| $0.98827 \leq X$ | 2 |

Table 1: Set One: Frequency table

B - Set Two: Frequency table

| Value Range | Count |
|----------------------------|-------|
| $1.00000 \leq X < 8375.00$ | 91 |
| $8375.00 \leq X < 16750.0$ | 10 |
| $16750.0 \leq X < 25125.4$ | 5 |
| $25125.4 \leq X < 33500.2$ | 5 |
| $33500.2 \leq X < 41875.0$ | 3 |
| $41875.0 \leq X < 50249.8$ | 2 |
| $50249.8 \leq X < 58624.6$ | 1 |
| $58624.6 \leq X < 66999.4$ | 0 |
| $66999.4 \leq X < 75374.2$ | 2 |
| $75374.2 \leq X$ | 1 |

Table 2: Set Two: Frequency table**IV - Code****A - Descriptive Analysis**

```
#AUTHOR: Joe Cloud
#PURPOSE: Perform simple descriptive analysis for Probability & Statistics for Engineers,
          project
#UTA FALL 2017

import numpy as np
import sys
import matplotlib.pyplot as plt
import tabulate

DATA_FILE = "../set_one/resistor_vals_offset.csv" # Set to default list
QUARTILES = [25, 50, 75]

if len(sys.argv) > 1:
    DATA_FILE = sys.argv[1]

OUTPUT_FILE = "results/" + DATA_FILE.split('/')[1].split('vals')[0]

def main():

    sample_vals = np.genfromtxt(DATA_FILE, delimiter=',')
    print(sample_vals)

    print("Min value is: %f" % min(sample_vals))
    print("Max value is: %f" % max(sample_vals))
```

```
sample_mean = np.mean(sample_vals)
print("Mean value is: %f" % sample_mean)

sample_std = np.std(sample_vals)
print("STD value is: %f" % sample_std)

# Calculate quartiles
sample_quarts = []
for quart in QUARTILES:
    sample_quarts.append(np.percentile(sample_vals, quart))

print("Quartiles: ", *sample_quarts, sep=', ')

generateTable(sample_vals)

# Construct box-and-whisker plot, a.k.a. boxplot
fig = plt.figure()
ax = plt.subplot(111)
ax.boxplot(sample_vals, 0, 'kp', 0)
fig.savefig(OUTPUT_FILE + 'boxplot.png', bbox_inches='tight')
fig.clf()

num_bins = 10

# Frequency table
frequency_table = np.histogram(sample_vals, bins=num_bins)
print(frequency_table)

# Histogram data
fig = plt.figure()
ax = plt.subplot(111)
ax.hist(sample_vals, bins=num_bins)
fig.savefig(OUTPUT_FILE + 'histogram.png', bbox_inches='tight')
fig.clf()

def generateTable(data):

    data_c = data.reshape(10, int(len(data)/10))
    print(data_c.shape)

    gen_table = tabulate.tabulate(data_c, tablefmt="latex")
    print(gen_table)

    outfile = open(OUTPUT_FILE + 'latex_gen.txt', 'w')
```



```
outfile.write("\n\n")
outfile.write("%s\n" % gen_table)

if __name__ == "__main__":
    main()
```

B - Data Processing: Set One

```
#!/usr/bin/env python3

# This simply offsets each value by 0.33

FILE_NAME="resistor_vals.csv"

def main():

    f = open(FILE_NAME, 'r')

    data = f.readlines()

    print(data)

    data_output = []
    for val in data:
        data_output.append(float(val) - 0.00033)

    outfile = open("resistor_vals_offset.csv", 'w')

    for val in data_output:
        outfile.write("%s\n" % val)

if __name__ == "__main__":
    main()
```

C - Data Processing: Set Two

```
#!/usr/bin/env python3

# This performs necessary pre-processing for raw data from 'last' command
```

```
FILE_NAME = "stripped_batch_three.log"

def main():

    f = open(FILE_NAME, 'r')
    data = f.readlines()
    data_convded = convToSeconds(data)
    print(data_convded)
    data_interval = calcInterval(data_convded)
    print(data_interval)

    outfile = open("logininterval_vals.txt", 'w')

    for val in data_interval:
        outfile.write("%s\n" % val)

def convToSeconds(data):

    in_seconds = []

    for login in data:
        # Takes value in format HH:MM:SS and tokenizes, adds up each delimited element
        # w/ respective multiplier.
        temp_tok = login.split(':')

        in_seconds.append(int(temp_tok[2])+int(temp_tok[1])*60+int(temp_tok[0])*3600)

    return in_seconds

def calcInterval(data):

    in_intervals = []
    # Go through and get the difference between T1 and T0 to determine login intervals
    # Must abs values bc for ex, T0 may be 23:34:46 and T1 00:12:49 where the difference would
    # be negative
    for i in range(0, len(data) - 1):
        in_intervals.append(abs(data[i] - data[i+1]))

    return in_intervals

if __name__ == "__main__":
    main()
```

```
#!/usr/bin/env bash
```

```
last -Fwx > last_batch_$1.log
```

```
cat last_batch_$1 | awk '{ print $7 }' > stripped_batch_$1.txt
```

V - Equipment Used



Figure 5: Set One: Fluke 289 w/ resistors

REFERENCES

Walpole, R.E., 2016. *Probability and Statistics for Engineers and Scientists*. Prentice Hall.